

Epistemic Semantics in Guarded String Models

Eric Campbell

Cornell University

ehc86@cornell.edu

Mats Rooth

Cornell University

mr249@cornell.edu

Abstract

Constructive and computable multi-agent epistemic possible worlds models are defined, where possible worlds models are guarded string models in an epistemic extension of Kleene Algebra with Tests. The account is framed as a formal language Epik (Epistemic KAT) for defining such models. The language is interpreted by translation into the finite state calculus, and alternatively by modeling propositions as lazy lists in Haskell. The syntax-semantics interface for a fragment of English is defined by a categorial grammar.

1 Introduction and Related Work

Linguistic semantics in the Montague tradition proceeds by assigning propositional *semantic values* to disambiguated sentences of a natural language. A proposition is a set or class of *possible worlds*. These are often assumed to be things with the same nature and complexity as the world we occupy (Lewis, 1986). But alternatively, one can work with small idealized models, in order to illustrate and test ideas. To build such models, spaces of worlds and individuals are stipulated as small finite sets, and semantic values of lexical items are constructed as functions or relations from these small sets. Such toy or idealized models are useful in research and in teaching, in that it is possible to represent propositions finitely and explicitly, and to calculate with them. The point of this paper is to scale up toy or idealized models to countable sets of worlds, and to constructive and computable modeling of epistemic alternatives for agents. We describe a certain systematic way of defining such models, and illustrate how to apply them in natural language semantics. The focus is on epistemic semantics and clausal embedding. The fundamental move is to identify possible worlds with strings of primitive events, so that propositions are sets

of strings. An advantage in this is that it allows for a mathematical description of an algebra of propositions, coupled with a computational representation using either lazy lists of strings, or finite state machines that describe sets of strings.

The approach taken here synthesizes five antecedents in a certain way. John McCarthy's *Situation Calculus* is the source of the idea of constructing possible worlds as event sequences (McCarthy, 1963; Reiter, 2001). The algebraic theory of *Kleene Algebra with Tests* characterizes algebras with elements corresponding to propositions and event types in our application (Kozen, 2001). The models we propose are an epistemic extension of guarded string models for KAT, where a unary operation interpreted as an existential epistemic modality is included for each agent. *Action models* in dynamic epistemic semantics introduced the technique of constructing epistemic models from primitive alternative relations on events, in order to capture the epistemic consequences of perceptual events (Baltag et al., 1999). This is the basis for our construction of epistemic alternative relations. Literature on finite state methods in linguistic semantics has used event strings and sets of event strings to theorize about tense and aspect in natural language semantics (Fernando, 2004, 2007; Carlson, 2009) and to express intensionality (Fernando, 2017). Literature on finite state intensional semantics has investigated how to do the semantics of intensional complementation, including indirect questions, in a setting where compositional semantics is expressed in a finite state calculus (Rooth, 2017; Collard, 2018). We adopt this in our syntax-semantics interface for English.

We begin with examples of event-sequence models. *The Elevator*. An elevator moves up and down in a four-story building, with floors numbered in the European fashion as 0,1,2,3. There are primitive events u (the elevator going up one floor), and

2001 or
2003?

d (the elevator going down on floor). In worlds v_1 and v_2 , the events shown in (1) transpire. The truth values for English sentences shown in (2) are observed.

(1)	v_1	u	it goes up from 0 to 1
		u	it goes up from 1 to 2
		d	it goes down from 2 to 1
		u	it goes up from 1 to 2

v_2	u	it goes up from 0 to 1
	u	it goes up from 1 to 2
	u	it goes up from 2 to 3

(2)	v_1	v_2	Sentence
	false	true	It's on floor 3.
	true	true	It has gone up.
	true	false	It has gone down.
	true	false	It could go up.

The Concealed Coin. Amy and Bob are seated at a table. There is a coin on the table under a cup, heads up. The coin could be heads or tails, and neither agent knows which it is. This initial situation is possible world w_1 . Two additional worlds w_2 and w_3 are defined by sequencing events after the initial state, with events interpreted as in (3). The truth values for English sentences shown in (5) are observed.

- (3) a_1 Amy peeks at heads, by tipping the cup. Bob sees she's peeking, but not what she sees.

b_1 Bob peeks at heads.

- (4) w_1

$$w_2 = w_1 a_1$$

$$w_3 = w_1 a_1 b_1$$

w_1	w_2	w_3	Sentence
false	true	true	Amy knows that it's heads.
false	false	true	Bob knows that it's heads.
false	false	true	Bob knows Amy knows it's heads.
false	true	true	Bob knows Amy knows whether it's heads or tails.

The events in the examples come with pre-conditions. The elevator can not go up if it is already on floor 3, so u has the pre-condition of the elevator being of floor 0, 1, or 2. Similarly d has the precondition that the elevator is on floor 1, 2

or 3. Amy can peek at heads only if the coin is heads up, so a_1 has the precondition of the coin being heads up. Let h be the Boolean proposition that the coin is heads up. In the other example, let q be the proposition that the elevator is on a high floor (2 or 3), and p be the proposition that it is on an odd floor (1 or 3). Then preconditions can be described by Boolean formulas, with h being the precondition of a_1 , and \bar{pq} being the precondition of u . Juxtaposition is used for Boolean conjunction, and the overbar for Boolean negation. Events come as well with a relation between prior and following state, for instance with u incrementing the floor. This is expressed using an operator “:” (read “and next”) that pairs Boolean formulas. The first line in (6) describes a_1 (Amy looking at heads) as happening only in an h state, and as not changing the state. Symmetrically, a_0 (Amy looking at tails) can happen only in a not- h state, and does not change the state. The third line says that u increments the floor, and can happen only on floors 0, 1, and 2. The fourth line describes d in similar terms. Plus is disjunction.

(6)	a_1	$h : h$
	a_0	$\bar{h}:\bar{h}$
	u	$(\bar{qp}):(\bar{qp}) + (\bar{qp}):(\bar{qp}) + (\bar{qp}):(\bar{qp})$
	d	$(\bar{qp}):(\bar{qp}) + (\bar{qp}):(\bar{qp}) + (qp):(\bar{qp})$

2 Epistemic guarded string models

In the discussion at the end of the last section, a sequence such as \bar{qp} can be viewed as a valuation of primitive test propositions, which is used to describe world state. The primitives are listed in fixed order, and left unmarked (indicating true) or marked with the overbar (indicating false). Suppose that in the coin example, we have an additional primitive stative proposition (or atomic test) t , interpreted as tails. Since a coin is heads or tails but not both, we want to allow the valuations ht and \bar{ht} , and disallow ht and \bar{ht} . This is enforced by a *state formula*, which is a simply Boolean formula, in this case (7a). Where B is a set of atomic tests and ϕ is a state constraint over B , \mathcal{A}_B^ϕ is the set of valuations of B that make formula ϕ true. The valuations are called atoms, because they correspond to the atoms of Boolean algebra of tests (Kozen, 2001).

Formulas like the ones in (6) that describe pre- and post-conditions are *effect formulas*. They are interpreted as defining relations between atoms, as defined in Figure 1. The atoms they relate are constrained by the state formula as well. For the

state formulas $(a \in B)$
 $\rho, \sigma, \varphi ::= a \mid 0 \mid 1 \mid \rho + \sigma \mid \rho\sigma \mid \bar{\rho}$

effect formulas

$\zeta, \eta ::= \rho : \sigma \mid \zeta + \eta \mid \zeta \& \eta \mid \bar{\zeta}$

$$\begin{aligned} \llbracket \rho : \sigma \rrbracket^\varphi &= \mathcal{A}_B^{\rho\varphi} \times \mathcal{A}_B^{\sigma\varphi} \\ \llbracket \zeta + \eta \rrbracket^\varphi &= \llbracket \zeta \rrbracket^\varphi \cup \llbracket \eta \rrbracket^\varphi \\ \llbracket \zeta \& \eta \rrbracket^\varphi &= \llbracket \zeta \rrbracket^\varphi \cap \llbracket \eta \rrbracket^\varphi \\ \llbracket \bar{\zeta} \rrbracket^\varphi &= \mathcal{A}_B^\varphi \times \mathcal{A}_B^\varphi \setminus \llbracket \zeta \rrbracket^\varphi \end{aligned}$$

Figure 1: Syntax of state formulas and syntax and semantics of effect formulas. Effect formulas denote relations between atoms. In a state formula, juxtaposition $\rho\sigma$ is conjunction.

heads-tails example, let the state formula and the effect formula for a_1 (Aly peeking at heads) be as specified in (7). Then \mathcal{A}_B^φ and the relation on atoms for the event a_1 are as given at the bottom in (7).

(7)	B	$\{h, t\}$
	state formula φ	$h\bar{t} + \bar{h}t$
	effect formula ζ for a_1	$h : h$
	\mathcal{A}_B^φ	$\{h\bar{t}, \bar{h}t\}$
	$\llbracket \zeta \rrbracket^\varphi$	$\{\langle h\bar{t}, \bar{h}t \rangle\}$

Figure 2 shows an Epik program that describes a possible worlds model for two agents with information about one coin, and events of the agents semi-privately looking at the coin. The line beginning with `state` enumerates B . The line beginning with `constraint` gives the state formula. The lines beginning with `event` declare events and their effect formulas. Finally the lines beginning with `agent` define *event alternative* relations for agents. Each clause with an arrow has a single event symbol on the left, and a disjunction of alternative events on the right of the arrow. The interpretation of Amy's alternatives for b_1 (Bill peeks at heads), is that when b_1 happens, for Amy either b_1 or b_0 (Bill peeks at tails) could be happening.

Kleene Algebra with Tests is an algebraic theory that is defined by equations and inequalities, which has model classes including guarded string models, relational models, finite models, and matrix models. This paper focuses on defining a concrete guarded string algebra, the elements of which are sets of guarded strings, from an Epik specification. Definitions and notation mostly follow (Kozen, 2001). Additional syntax and semantics is

```
state h t
constraint h!t + t!h
event a1 h:h
event a0 t:t
event b1 h:h
event b0 t:t
agent aly
  a1 -> a1
  a0 -> a0
  b1 -> b1 + b0
  b0 -> b1 + b0
agent bob
  b1 -> b1
  b0 -> b0
  a1 -> a1 + a0
  a0 -> a1 + a0
```

Figure 2: Epik program describing a possible-worlds event sequence model for two agents with information about one coin, and events of the agents semi-privately looking at the coin.

included to cover multi-agent epistemic semantics. Guarded strings over a finite alphabet P are like ordinary strings, but with atoms alternating with the symbols from P . In the algebra described by Figure 2, is the set of events $\{a_1, a_0, b_1, b_0\}$, and in the elevator example, P is $\{u, d\}$. In the elevator example, B is $\{p, q\}$, and in the coin example it is $\{h, t\}$. Assuming a trivially true state formula ρ , \mathcal{A}_B^ρ in the elevator example is $\{\bar{q}\bar{p}, \bar{q}p, q\bar{p}, qp\}$, which we write $\{\hat{0}, \hat{1}, \hat{2}, \hat{3}\}$. In the coin example, as we already saw in (7), \mathcal{A}_B^φ is $\{h\bar{t}, \bar{h}t\}$, for which we use the shorthand $\{H, T\}$. Guarded strings are strings of events, alternating with atoms, and beginning and ending with atoms. (8) gives the encoding as guarded strings of the worlds in (1) and (3). The length of a guarded string p , written $|p|$ is the number of events in p . An atom such as H is a guarded string of length 0.

(8)	World	Guarded string	Length
v_1	$\hat{0} u \hat{1} u \hat{2} d \hat{1} u \hat{2}$	4	
v_2	$\hat{0} u \hat{1} u \hat{2} u \hat{3}$	3	
w_1	H	0	
w_2	$H a_0 H$	1	
w_3	$T a_0 T b_0 T$	2	

The discussion of (4) mentioned building worlds by incrementing worlds with events. This is accomplished in guarded string models with fusion

product \diamond , a partial operation that combines two guarded strings, subject to the condition that the atom at the end of the first argument is identical to the atom at the start of the second one. (9) gives some examples.

$$(9) \begin{aligned} \hat{0} u \hat{1} u \hat{2} d \hat{1} \diamond \hat{1} u \hat{2} &= \hat{0} u \hat{1} u \hat{2} d \hat{1} u \hat{2} \\ \hat{0} u \hat{1} u \hat{2} d \hat{1} \diamond \hat{2} u \hat{3} &= \text{undefined} \\ H \diamond H a_1 H &= H a_0 H \\ H \diamond T a_1 T &= \text{undefined} \end{aligned}$$

Rather than individual guarded strings, elements of a guarded string model for KAT are sets of guarded strings. In the application, these elements have the interpretation of propositions (sets of possible worlds). In a free guarded string model for KAT, any event can be adjacent to any atom in a guarded string that is an element of the underlying set for the algebra. We instead impose the constraints coming from the state and effect formulas. (10) defines the well-formed guarded strings determined by an Epic specification. (i) says that each atom is consistent with the state constraint, and (ii) says that each constituent token event $\alpha_i e_i \alpha_{i+1}$ is consistent with the effect constraint on e_i .¹

- (10) Given P , B , a state formula φ , and an effect formula ζ_e for each event e in P , $\alpha_0 e_0 \dots e_n \alpha_{n+1}$ is well-formed iff
- (i) $\alpha_i \in \mathcal{A}_B^\varphi$ ($0 \leq i \leq n$), and
 - (ii) $\langle \alpha_i, \alpha_{i+1} \rangle \in [\![\zeta_e]\!]^\varphi$, ($0 \leq i \leq n$).

Well-formed guarded strings have the interpretation of worlds in the application to natural-language semantics. The set of possible worlds in the Kripke frame determined by an Epic specification is the set of well-formed guarded strings. At this point, we could say in the way familiar from the type theory of possible worlds semantics that any set of worlds is a proposition, so that the set of propositions is the power set of the set of worlds (??). We will instead define a more restrictive set of propositions corresponding to the regular sets of strings. This is deferred to the next section.

In the natural language application, sets of well-

¹An alternative is to define equations such as $\bar{\phi} = 0$ (from the state formula ϕ) and $a_1 = h a_1 h$ (from the effect formula $h : h$ for event a_1), and construct a quotient algebra from the equivalence relation generated by these equations. This results in equating sets of guarded strings in the free algebra that differ by guarded strings that are ill-formed according to the state and effect formulas. In the development in the text, we instead use a set of guarded strings that are well-formed according to the state and effect formulas as the representative of the equivalence class.

formed guarded strings have the additional interpretation of event types. An event-type is something that can “happen” in different worlds. The event of the elevator going up is modeled not as the bare event symbol u or its unit set $\{u\}$, but as the set of guarded strings $\hat{0} u \hat{1}, \hat{1} u \hat{2}, \hat{2} u \hat{3}$. The event of the elevator going up in a given world w corresponds to w being incremented to form $w \diamond x$, where x is an element of the event type. (11) gives the event types for the examples.

(11)		Event type
u		$\{\hat{0} u \hat{1}, \hat{1} u \hat{2}, \hat{2} u \hat{3}\}$
d		$\{\hat{1} d \hat{0}, \hat{2} d \hat{1}, \hat{3} d \hat{2}\}$
a_1		$\{H a_1 H\}$
a_0		$\{T a_0 T\}$
b_1		$\{H b_1 H\}$
b_0		$\{T b_0 T\}$

The construction so far defines a set of worlds W from an Epic specification. Normally and in our examples, W is countably infinite, though some choices of effect formulas can result in a finite set of worlds. The next step is to define an alternative relation R_a on worlds for each agent a . This will result in a Kripke frame $\langle W, R_a, \dots \rangle$ in the standard sense, consisting of a set of worlds, and a world-alternative relation for each agent (??). An Epic specification defines an alternative relation on bare events for each agent a , which we denote as \hat{R}_a . This should be lifted to a relation R_a on worlds. The rough idea is that when a world w is incremented with an event e , in the resulting world $w \diamond e$, epistemic alternatives for agent a are of the form $w' \diamond e'$, where w' is an alternative to w for a in w , and e' is an event-alternative for a (Moore –, Baltag –).² This need to be implemented in a way that takes account of pre- and post-conditions for events. For this, our approach is to refer the definition of well-formed guarded strings. (12) defines a relation on worlds from a relation on bare events.

²In this it is important that the event-alternative relation for an agent is constant across worlds. We anticipate that the definition given here produces results equivalent to what is found in literature on event alternatives in dynamic epistemic semantics, though we have not verified this. That literature primarily focuses on mapping an epistemic model for a single time and situation to another, and uses general first-order models, rather than guarded string models. See —, and articles in —. We picked up the idea from papers by Moss and his colleagues, together with Moore () and subsequent literature in situation theory. This previous literature is motivated by epistemic logic and AI planning, rather than natural language semantics as in our application.

- (12) Let a set of guarded strings W over events P and primitive tests B , and a relation \hat{R} on P be given. The corresponding relation R holds between a guarded string

$$\alpha_0 e_0 \dots e_n \alpha_{n+1}$$

in W and a guarded string q iff q is an element of W and is of the form

$$\alpha'_0 e'_0 \dots e'_n \alpha'_{n+1},$$

where for $0 \leq n$, $\langle e_i, e'_i \rangle \in \hat{R}$.

This requires that in the alternative world q , each constituent event e'_i is an alternative to the event e_i in the base world according to \hat{R} . Compatibilities between events in the alternative world are enforced by the requirement that the alternative q is an element of W , so that state and effect formulas are enforced. Consider a scenario like the one from Figure 1, with an additional agent Cem. The base world Tb_0c_0 is one where the coin is tails, and first Bob looks at tails, and then Cem looks at tails. The first event b_0 has the alternatives b_0 and b_1 for Aly, and the second event c_0 has the alternatives c_0 and c_1 for Aly. This results in the combinations b_0c_0 , b_0c_1 , b_1c_0 , and b_1c_1 . But these are filtered by post- and pre-conditions of events in the alternative world, so that the set of alternatives for Aly in Tb_0c_0 is $\{Tb_0c_0, Hb_1c_1\}$, with two world-alternatives instead of four.

3 The language of Epistemic KAT

In addition to a set of possible worlds and an algebra of propositions, the application in natural language requires a logical language for naming propositions. The standard language for Kleene algebra with tests has the signature $\langle K, +, \cdot, *, \bar{\cdot}, 0, 1 \rangle$ (Kozen, 2001). In a guarded string model for KAT, K is a set of guarded strings, $+$ is set union, \cdot is fusion product raised to sets, $*$ is Kleene star, $\bar{\cdot}$ is complement for tests, 0 the empty set (functioning as an identity for 0 , and as False for tests) and 1 is the set of atoms (functioning as the identity element for \cdot , and as True for tests). To this signature we add a unary modal operation \diamond_a for each agent, and a unary complement operation $\bar{\cdot}$ on elements of K . The latter is included because natural languages have sentence negation, and it should be possible to represent the propositional complement of a formula that captures the semantics of a natural language sentence. In addition, universal

state terms σ as in Figure 1

events $e \in P$

$$p, q ::= e \mid \sigma \mid p + q \mid pq \mid p^* \mid \neg p \mid \diamond_a p$$

Derived Operators

$$\square_a p \triangleq \neg \diamond_a \neg p \quad p \& q \triangleq \neg (\neg p + \neg q) \quad - \triangleq \sum P$$

Figure 3: The language of Epik terms and key derived operators

box modalities are defined as duals of existential diamond modalities.

With modalities and propositional negation added, the signature of n -agent epistemic KAT is $\langle K, +, \cdot, *, \bar{\cdot}, 0, 1, \neg, \diamond_1 \dots \diamond_n \rangle$. Figure 3 defines the syntax of the language, but using juxtaposition in place of the product symbol \cdot . As illustrated in (13), terms in this language are used to represent the propositional semantic values of English sentences.³

- (13) Bob knows that Amy knows whether it's tails.
 $\square_b(\square_a(\neg 0)t + \square_a(\neg 0)\neg t)$

Each term in the language is interpreted as a set of guarded strings. Where p is a term, $\llbracket p \rrbracket^{B, P, \varphi, \zeta}$ is the interpretation of the term. The superscript captures dependence on an Epik specification.

interprets a term p as a set of guarded strings consistent with state constraints φ and effect relations E^φ . A bare event e is interpreted as the set $\mathcal{E}^\varphi(e)$. The constant 0 is interpreted as the empty set. The constant 1 is interpreted as the set A_B^φ , i.e. $\{ht, \bar{ht}\}$ in the coin example. The operation $p + q$ is interpreted as the union of the interpretations of p and q . The operation $p; q$ interpreted as the fusion product of the interpretations of p and q raised to sets: $X \diamond Y = \{x \diamond y \mid x \in X, y \in Y, x \diamond y \text{ defined}\}$. The operation p^* is Kleene star, which is interpreted to be least upper bound of the iterated fusion product of the denotation of p with itself. Subsets of 1 indicated by propositions ψ are also elements of K , and these form the Boolean algebra of tests, denoting all atoms admitted by the constraint φ and the boolean test ψ .

The unary epistemic alternative operation \diamond_a for each agent a , is interpreted using Kripke semantics, as pre-image relative to a fixed relation \hat{R}_a between guarded strings, $\diamond_a p = \{u \mid \exists v. v \in \llbracket p \rrbracket^{\phi, E} \wedge u \hat{R}_a v\}$. Here u and v are guarded strings, while x is an element of K .

³See Sections 6 and 7 for explanation and further examples.

$$\begin{aligned}
\llbracket 0 \rrbracket^{B,P,\varphi,\zeta} &\triangleq \emptyset \\
\llbracket 1 \rrbracket^{\varphi,\mathcal{E}} &\triangleq \mathcal{A}_B^\varphi \\
\llbracket e \rrbracket^{\varphi,\mathcal{E}} &\triangleq \hat{\mathcal{E}}^\varphi(e) \\
\llbracket \psi \rrbracket^{\varphi,\mathcal{E}} &\triangleq \mathcal{A}^\varphi \psi \\
\llbracket p + q \rrbracket^{\varphi,\mathcal{E}} &\triangleq \llbracket p \rrbracket^{\varphi,\mathcal{E}} \cup \llbracket q \rrbracket^{\varphi,\mathcal{E}} \\
\llbracket p; q \rrbracket^{\varphi,\mathcal{E}} &\triangleq \llbracket p \rrbracket^{\varphi,\mathcal{E}} \diamond \llbracket q \rrbracket^{\varphi,\mathcal{E}} \\
\llbracket p^* \rrbracket^{\varphi,\mathcal{E}} &\triangleq \bigcup_n \underbrace{\llbracket p \rrbracket^{\varphi,\mathcal{E}} \diamond \dots \diamond \llbracket p \rrbracket^{\varphi,\mathcal{E}}} \\
\llbracket \neg p \rrbracket^{\varphi,\mathcal{E}} &\triangleq \llbracket _^* \rrbracket^{\varphi,\mathcal{E}} \setminus \llbracket p \rrbracket^{\varphi,\mathcal{E}} \\
\llbracket \Diamond_a p \rrbracket &\triangleq \{x \mid x \hat{R}_a y, y \in \llbracket p \rrbracket^{\varphi,\mathcal{E}}\}
\end{aligned}$$

Figure 4: Interpretation of Epik terms as sets of guarded strings

$$\begin{aligned}
\llbracket 0 \rrbracket^{\varphi,\mathcal{E}} &\triangleq \{\} \\
\llbracket 1 \rrbracket^{\varphi,\mathcal{E}} &\triangleq \mathcal{A}_B^\varphi \\
\llbracket e \rrbracket^{\varphi,\mathcal{E}} &\triangleq \hat{\mathcal{E}}^\varphi(e) \\
\llbracket \psi \rrbracket^{\varphi,\mathcal{E}} &\triangleq \mathcal{A}^\varphi \psi \\
\llbracket p + q \rrbracket^{\varphi,\mathcal{E}} &\triangleq \llbracket p \rrbracket^{\varphi,\mathcal{E}} \cup \llbracket q \rrbracket^{\varphi,\mathcal{E}} \\
\llbracket p; q \rrbracket^{\varphi,\mathcal{E}} &\triangleq \llbracket p \rrbracket^{\varphi,\mathcal{E}} \diamond \llbracket q \rrbracket^{\varphi,\mathcal{E}} \\
\llbracket p^* \rrbracket^{\varphi,\mathcal{E}} &\triangleq \bigcup_n \underbrace{\llbracket p \rrbracket^{\varphi,\mathcal{E}} \diamond \dots \diamond \llbracket p \rrbracket^{\varphi,\mathcal{E}}} \\
\llbracket \neg p \rrbracket^{\varphi,\mathcal{E}} &\triangleq \llbracket _^* \rrbracket^{\varphi,\mathcal{E}} \setminus \llbracket p \rrbracket^{\varphi,\mathcal{E}} \\
\llbracket \Diamond_a p \rrbracket &\triangleq \{x \mid x \hat{R}_a y, y \in \llbracket p \rrbracket^{\varphi,\mathcal{E}}\}
\end{aligned}$$

Figure 5: Interpretation of Epik terms as sets of guarded strings

The complement operation p^c is complement at the level of sets of guarded strings, defined to every set the universe $_^*$ except those denoted by p . Note that $\llbracket 0^c \rrbracket^{\varphi,\mathcal{E}} = \llbracket _^* \rrbracket^{\varphi,\mathcal{E}}$.

Summing up, given an Epik program with n agents, we construct state constraints φ , effect relations \mathcal{E} , and a concrete guarded string model for each term using $\llbracket - \rrbracket^{\varphi,\mathcal{E}}$. 0^c is the set of worlds, and it may be countably infinite. \Diamond_{a_i} is an epistemic modality for the i th agent. Or referring to the Kripke relations \hat{R}_i , the construction defines a multi-agent Kripke frame $\langle 0^c, \hat{R}_1, \dots, \hat{R}_n \rangle$ (usually a countable one) from an Epik specification. The frame consists of a set of worlds, and an epistemic-alternative relation for each agent. These models are used as a target for natural-language interpretation in Section 5 and Section 6, where we obtain semantic values such as $\llbracket \text{Amy knows that it's heads} \rrbracket$ and $\llbracket \text{Bob knows that Amy knows whether it is heads or tails, and does not know that it's heads} \rrbracket$ as elements of K . Concretely the propositions are sets of guarded strings (usually countable ones), construed as sets of worlds as they figure in possible worlds semantics for natural language.

4 Translation into the finite state calculus

The finite state calculus is an algebra of regular sets of strings and regular relations between strings that was designed for use in computational phonology and morphographemics (Kaplan and Kay, 1994; Beesley and Karttunen, 2003). Current implementations allow for the definition of functions with the status of defined operators on regular sets and relations (Hulden, 2009; Karttunen, 2010). Such definitions are used to define an embedding of epistemic KAT in a string algebra. The methodology follows Section 2 closely. Let \mathcal{K} be an epistemic algebra as described in Section 2. A given element of \mathcal{K} is represented in the string algebra by the very same set of strings, i.e. by a set of strings that have the form of a sequence of bare event symbols, with interleaved Boolean vectors. Product in the KAT can not be modeled as concatenation in the string algebra, because this would not enforce identity of states, and would result in lengthening Boolean vectors at the concatenation point. Instead, KAT product and KAT Kleene star are defined operations in the string algebra, see Figure 6. The operations concatenate in the string algebra, delete strings with non-matching tests, and then delete the second of two tests create

```

St Tests such as 0 1 1 0. The length is the number
of generators.
UnequalStPair Sequence of two unequal
tests such as 0 1 1 0 0 1 1 1, differing in one
or more positions.
define Wf0 ~[$ UnequalStPair];
String that doesn't contain a non-matching test
pair.
define Squash St -> 0 || St _;
Rewrite relation deleting the second of two tests.
define Cn(X, Y)
  [[X Y] & Wf0] .o. Squash].l;
KAT product in Fst, where & is intersection, .o.
is relation composition, and .l is relation image.
define Kpl(X)
  [[[X+] & Wf0] .o. Squash].l;
define Kst(X) St | Kpl(X);
KAT Kleene plus and Kleene star in Fst. The Fst
operation | is union.

```

Figure 6: Translation into Fst of KAT product and KAT Kleene star.

a well-formed guarded string.

A given bare event such as a_1 (Aly looks at heads) is in the KAT algebra a set of bare events decorated with compatible tests on each side, semantically $\{10a_110\}$ in this case. This is a unit set rather than a guarded string, because elements of the KAT algebra are sets. Worlds in the KAT algebra are defined by sequencing events using Kst . The operation enforces compatibility of states, so that $(a_1 + a_0)(b_1 + b_0)$ contains two worlds rather than four. The program in Figure 2 as interpreted in FST defines a countably infinite set of possible worlds by KAT Kleene closure as $Kst(a_1 + a_0 + b_1 + b_0)$, and an algebra of propositions as regular sets of strings drawn from this space of worlds.

It remains to define an epistemic alternative relation on worlds for each agent. The relevant information in Figure 2 is a relation between bare events for each agent. This determines a relation in the guarded string algebra a relation between bare events decorated with compatible tests. For agent Aly, this is the relation described in (14) as a set of ordered pairs.

$$(14) \quad \left\{ \begin{array}{l} \langle 10a_110, 10a_110 \rangle, \\ \langle 01a_001, 01a_001 \rangle, \\ \langle 10b_110, 10b_110 \rangle, \\ \langle 10b_110, 01b_001 \rangle, \\ \langle 01b_001, 10b_110 \rangle, \\ \langle 01b_001, 01b_001 \rangle \end{array} \right\}$$

```

define RelKpl(R) Squash.i.o.
Wf0.o. [R+] .o. Wf0 .o. Squash
b a b c
a Relational Kleene plus in the string algebra
b Constrain domain and co-domain to contain
no unmatched tests.
c Reduce doubled tests to a single
test in the domain and co-domain.
Squash.i is the inverse
of Squash.
define Kst(R) [St.x.St] | Kpl(X);
The Fst operation .x. is Cartesian product.

```

Figure 7: Definition in Fst of the Kleene concatenation closure of a relation between guarded strings.

The relation on decorated events needs to be generalized to a relation of worlds. The principle for this is that an epistemic alternative to a world of the form we is a world of the form vd , where v is a world-alternative to w , d is an event-alternative to e , and vd is defined (i.e. the world alternative v satisfies the pre-conditions of the event alternative d). This principle is found in earlier literature (Moore 198x, Baltag, Moss and Solecki 20xx). In the construction in Fst, the definition of world alternatives takes a simple form. Where R_a is the relation on decorated events for agent a , the corresponding relation on worlds in is the Kleene closure of R_a . Where R and S are relations, the concatenation product of R and S is the set of pairs of the form $\langle x_1x_2, y_1y_2 \rangle$, where $\langle x_1, y_1 \rangle$ is in relation R , and $\langle x_2, y_2 \rangle$ is in relation S . The Kleene closure of relation R is $\cup_{n \geq 0} R^n$, where R^n is the n -times concatenation product of R with itself (the 0-times concatenation product is $\llbracket 1 \rrbracket^{\varphi, \mathcal{E}}$). This is an operation in the finite state calculus. Figure ?? defines the corresponding operation in the guarded string algebra. The epistemic alternative relation on worlds for an agent is then defined as the concatenation closure of the event alternative relation for the agent.

Other operations in the guarded string algebra as defined in FST are simpler. Union is union in the string algebra. The complement of a proposition is complement relative to the set of worlds, as defined by set difference in the string algebra.

This scheme provides for an interpretation of the language of Epik terms that is defined in Figure 3. An Epik specifications such as Figure 1 is translated to a straight-line program of the finite

state calculus that defines constants and functions, including the ones from Figures 2 and 3. In Xfst or Foma, which are interpreters for the finite state calculus, the program is read, and then propositional terms can be mapped to finite state machines that represent sets of guarded strings, interpreted as propositions.

5 Bounded Lazy Interpretation

We also implement the semantics of Epik terms using lazy lists (?) in Haskell, rather than the direct interpretation as sets. Unfortunately, regular expressions, and hence also Epik, can denote infinite sets of strings, for example the term $_^*$. Normally regular languages are represented using a finite coalgebra (). However the non-distributivity of \diamond accross ; complicates this construction⁴. To sidestep this, we parameterize the interpretation function on a positive integer n and only produce guarded strings of length n or less.

We translate the Epik terms into a Haskell algebraic datatype that represents terms with the same signature as described in Section 2. We then parameterize the interpretation function, $LpM_n^{\varphi, \mathcal{E}}$ on an integer n that describes the maximum length of a string we will produce.

The bounded interpretation into lists of strings is very similar to the unbounded interpretation into sets of strings, except for the bounds checking. The full details are shown in Figure ???. First note that when $n = 0$, we simply return the empty list, denoted $[]$. Terms of the form 0 , 1 , e , and ψ have the same denotation as before, translated into a list (for a set X , $[X]$ is a list with the same elements as X arbitrarily ordered). The term $p + q$ denotes the list concatenation (written $++$) of the denotation of p and of q . The term $p; q$ denotes the fusion product (lifted to lists) of the denotations of p and in q restricted to only those strings shorter than n (for a list l , $l|_n$ filters out elements longer than n). The denotation of p^* uses the fact that p^* and $1 + p; p^*$ are equivalent, and decrements the size threshold on the recursive denotation of p^* by i , where i is the length of the longest (nonzero) string in the denotation of p , making sure to filter out guarded strings that are too long. The denotation of $\neg p$ is the strings that occur in $_^*$ and not in p (the \ operator on lists is analogous to the \ operator on sets). The denotation of $\diamond_a p$ is analogous to the set

⁴Axiomatic and Coalgebraic models for Epik are open questions

$$\begin{aligned}
LpM_0^{\varphi, \mathcal{E}} &\triangleq [] \\
L0M_n^{\varphi, \mathcal{E}} &\triangleq [] \\
L1M_n^{\varphi, \mathcal{E}} &\triangleq [\mathcal{A}_B^\varphi] \\
LeM_n^{\varphi, \mathcal{E}} &\triangleq [\hat{\mathcal{E}}^\varphi(e)] \\
L\psi M_n^{\varphi, \mathcal{E}} &\triangleq [\mathcal{A}_B^{\varphi \psi}] \\
Lp + qM_n^{\varphi, \mathcal{E}} &\triangleq LpM_n^{\varphi, \mathcal{E}} ++ LqM_n^{\varphi, \mathcal{E}} \\
Lp; qM_n^{\varphi, \mathcal{E}} &\triangleq (LpM_n^{\varphi, \mathcal{E}} \diamond LqM_n^{\varphi, \mathcal{E}})|_n \\
Lp^*M_n^{\varphi, \mathcal{E}} &\triangleq [] + (LpM_n^{\varphi, \mathcal{E}} \diamond Lp^*M_m^{\varphi, \mathcal{E}})|_n \\
\text{where } i &= \max\{1, \min\{|g| \mid g \in LpM_n^{\varphi, \mathcal{E}}\}\} \\
m &= \max\{0, n - i\} \\
L_{\neg}pM_n^{\varphi, \mathcal{E}} &\triangleq L_{\neg}^*M_n^{\varphi, \mathcal{E}} \setminus LpM_n^{\varphi, \mathcal{E}} \\
L\diamond_a pM_n^{\varphi, \mathcal{E}} &\triangleq [g' \mid g' \hat{R}_a g, \text{ for } g \text{ in } LpM_n^{\varphi, \mathcal{E}}]
\end{aligned}$$

Figure 8: Bounded interpretation using lazy lists

denotation, depicted in Figure ??-comprehension notation⁵.

To convert from a list of guarded strings l to a set of guarded strings, we simply write $[l]$. Note that for any p , φ , \mathcal{E} , and n , $[LpM_n^{\varphi, \mathcal{E}}] = \{g \mid g \in [p]^{\varphi, \mathcal{E}}, |g| \leq n\}$.

The lists we use here are *lazy* (as opposed to *strict*), which broadly means that computation is delayed until the value is needed. This allows us to avoid computing large, unnecessary iterations. In the elevator example, this means that $\hat{0}; u; u; u^*$ will unroll the u^* 0 times, once, and then discover that it is impossible to unroll it a second time, because $L\hat{3}; uM_n^{\varphi, \mathcal{E}} = []$, for any n . So $L\hat{3}; uM_n^{\varphi, \mathcal{E}} \diamond Lu^*M_{n-1}^{\varphi, \mathcal{E}}$ will return $[]$ without having to compute the full fixpoint of $LuM^{\varphi, \mathcal{E}}$. Similar behavior occurs once $n = 0$.

Conversely, if we used sets (as in the math) instead of lists we would need to constantly verify the set invariant (that every element is unique) which means processing every element in the set. Verifying that every element in $Lu^*M^{\varphi, \mathcal{E}}$ is unique would require the full computation of the fixpoint.

6 Syntax-semantics interface

An architecture of interpretation by translation is employed, where English sentences are mapped to terms in the logical language (-) via an interpreted grammar, and these terms are in turn interpreted

⁵List comprehension notation is analogous to set builder notation, except that it is written using square brackets. The order preservation is indicated by the keyword `for`

Basic statives	It's heads. It's tails.
That-complement	Amy knows that it's heads.
Wh-complement	Amy knows whether its heads.
Negation	Bob doesn't know that it isn't heads.
Tensed and base verbal forms	Bob knows that it's heads. Bob doesn't know that it's heads.
Sentence conjunction	It's heads and Bob doesn't know that it's heads
Predicate conjunction	Bob knows that Amy knows whether it's heads and doesn't know that Amy knows that it's heads.

Figure 9: Phenomena covered in the English grammar fragment.

as propositions (sets of possible worlds). For the latter, there are options of translation into the finite state calculus in order to represent propositions as finite state machines (Section 3), and representation in Haskell via lazy lists of guarded strings (Section 4). The grammar is a semantically interpreted multi-modal categorial grammar, consisting of a lexicon of words, their categorial types, and interpretations in a logical lambda language. Figure 9 lists phenomena that are covered.

As illustrated towards the end, there is recursion through conjunction and verbal complementation, so that the language is infinite, and includes talk of beliefs about beliefs, or in general, talk of arbitrarily iterated belief.

Figure 10 gives illustrative lexical entries. The grammar and semantics are in certain way optimized for a simple fragment of English concerned with clausal complementation. The agent names *Amy* and *Bob* contribute the epistemic alternative relations for those agents, rather than individuals. This is possible because the agents are never arguments of extensional predicates, so what matters about the agents is their epistemic alternative relations. The root verb *know* contributes existential modal force. The complementizers *that* and

Amy	e	R_a
Bob	e	R_b
it	d	d
heads	$d \setminus_D t$	$\lambda x.0^c.h$
tails	$d \setminus_D t$	$\lambda x.0^c.!h$
is	$(d \setminus t) / (d \setminus_D t)$	$\lambda P.\lambda x.Px$
knows	$(e \setminus t) / M t$	$\lambda p.\lambda R.\Diamond Rp$
that	$((e \setminus t) / M t) \setminus (e \setminus t) / t$	$\lambda p.\lambda m.\lambda R.\sim(m(\sim p)R)$
whether	$((e \setminus t) / M t) \setminus (e \setminus t) / t$	$\lambda p.\lambda m.\lambda R.\sim(m(\sim p)R) + \sim(mpR)$

Figure 10: Partial categorial grammar lexicon. The first column has a word form. the second column a categorial type, and third column a semantic translation in a logical language that extends the Epik term language with lambda.

whether are the heads of their dominating clauses, and assemble an alternative relation, modal force, and proposition contributed by the complement. These complementizers introduce the dual via two negations, in order to arrive universal modal force. These moves are offered here as a way of constructing a compact interpreted grammar. They can easily be reformulated in a more comprehensive interpreted grammar of English.

Multimodal categories such as \setminus_D and \setminus_M are used to control the derivation. For instance the category of *heads* $d \setminus_D t$. The dummy expletive subject *it* has category d , but the phrase *it heads* of category t can not be formed, because \setminus_D is not syntactically active as a function. Instead *it is heads* can be formed with a predicator *is* of category $(d \setminus t) / (d \setminus_D t)$. (This uses Lambek/Bar-Hillel notation for slashes, so that $(d \setminus t) / (d \setminus_D t)$ combines with $d \setminus_D t$ on the right to give a value that combines with d on the left to give t .) Similarly *knows* has a category with the top-level slash $/ M$, and combines to form a sentence as an argument of *that* or *whether*, which has a category that looks for the category of *know* on the left, after combining with a complement sentence on the right.

The semantic translations in the third column of Figure 10 use the Epik term language, incremented with lambda. The body of $\lambda x.0^c.h$, which is the semantic lexical entry for *heads*, is a term denoting the set of all worlds where the coin is heads, represented as the set of all guarded strings that end with a Boolean valuation where the primitive proposition *h* (it's heads) is true. There is λx at

the front because of a correspondence the grammar formalism uses a correspondence between syntactic and semantic types. However, it does not bind anything, because sentences such as *it isn't heads* have an expletive subject. The body $\Diamond R p$ of $\lambda p. \lambda R. \Diamond R p$, which is the semantic lexical entry of *knows*, is an Epik term denoting the pre-image of the set of worlds p according to the relation \hat{R} between guarded strings that is determined by the event-level relation R . This is not the right semantics for *Amy knows that it's heads*, because it is an existential modality $\Diamond R p$, rather than an universal modality $\Box R p$. This is corrected by the complementizer *that* or *whether*, which introduces the dual.

Sentences are parsed with a chart parser for categorial grammar. The semantics for complex phrases are obtained by syntactic application of semantic translations, accompanied by beta reduction. Semantic terms in the parsing formalism are expressions of untyped lambda calculus. The grammar is set up so that lambda is eliminated by beta reduction in the semantic term corresponding to a sentence. In consequence, the semantic term translating a sentence is a term of the Epik language ($-$). Such a term designates a set of possible words (guarded strings) in the possible worlds model determined by an Epik specification such as the one in Figure 1. (15a) is an English sentence with predicate conjunction and three levels of clausal embedding. Using the grammar and parser, the sentence is mapped to the Epik term (??). Using the result from Section 3, this term can be mapped in an implementation of the finite state calculus to a finite state machine that represents a countably infinite set of possible worlds, represented as guarded strings. Using the result from Section 4, it can be mapped to an infinite lazy list of guarded strings, representing the same set of possible worlds. Either of these is a concrete computational representation of the propositional semantic value $\llbracket \text{Amy knows that Bob knows that Amy knows whether it is heads and knows that Bob does not know that Amy knows that it is tails} \rrbracket^o$, in the familiar sense of Montague semantics for natural language.

- (15) a. Amy knows that Bob knows that Amy knows whether it is heads and knows that Bob does not know that Amy knows that it is tails.

b.

7 Examples and discussion

Page breakdown

1,2	3.75	
3	1.25	FST translation
4	1.25	Haskell Epik
5	0.75	English CG fragment
6	1.0	Examples and discussion

This stuff was at the end of Section 2, it became redundant. Also it refers to the interpretation function, which is introduced in Section 3.

$$(16) \hat{R} \triangleq \{ \langle x, y \rangle \mid \exists u_1 \dots \exists u_n \exists v_1 \dots \exists v_n \wedge \\ u_1 R v_1 \wedge \dots \wedge u_n R v_n \wedge \\ x \in \llbracket u_1; \dots; u_n \rrbracket^{\varphi, \mathcal{E}} \wedge \\ y \in \llbracket v_1; \dots; v_n \rrbracket^{\varphi, \mathcal{E}} \}$$

This defines an epistemic alternative to world x to be a world of the same length, where each component event in the alternative is an event-alternative to the event in corresponding position in the base world. Fusion product enforces preconditions and a correspondence between pre-states and post-states of events on both sides of the epistemic alternative relation. This provides for finitely specifiable construction of epistemic models that reflect intuitions about information exchange and epistemic consequences of perceptual events. See Section 6 for linguistic examples. Since an epistemic alternative has the same length as its base world, it follows from the construction that agents know how many events have transpired in their base worlds.

Fortunately, Epik provides a syntax to specify the input-output behavior via the `event` declarations. From these we construct a function \mathcal{E} from agents to relations on atoms that specify the effects each action can have. Notably for every declaration event $e \zeta$, we have $\mathcal{E}^\varphi(e) = \llbracket \zeta \rrbracket^\varphi$. We can lift \mathcal{E}^φ to $\hat{\mathcal{E}}$ which decorates the events with their permitted atoms, that is $\hat{\mathcal{E}}^\varphi(e) = \{ \alpha \ e \ \beta \mid \alpha, \beta \in \mathcal{E}^\varphi(e) \}$. See the examples in (11).

References

- Balag, A., L. S. Moss, and S. Solecki (1999). The logic of public announcements, common knowledge, and private suspicions.
- Beesley, K. R. and L. Karttunen (2003). *Finite State Morphology*. Center for the Study of Language and Inf.
- Carlson, L. (2009). *Tense, Mood, Aspect, Diathesis*. Book ms., University of Helsinki.
- Collard, J. (2018). Finite state reasoning for presupposition satisfaction. In *Proceedings of the First International Workshop on Language Cognition and Computational Models*, pp. 53–62.
- Fernando, T. (2004). A finite-state approach to events in natural language semantics. *Journal of Logic and Computation* 14(1), 79–92.
- Fernando, T. (2007). Observing events and situations in time. *Linguistics and Philosophy* 30(5), 527–550.
- Fernando, T. (2017). Intensions, types and finite-state truthmaking. In *Modern Perspectives in Type-Theoretical Semantics*, pp. 223–243. Springer.
- Hulden, M. (2009). Foma: a finite-state compiler and library. In *Proceedings of the 12th Conference of the European Chapter of the Association for Computational Linguistics: Demonstrations Session*, pp. 29–32. Association for Computational Linguistics.
- Kaplan, R. M. and M. Kay (1994). Regular models of phonological rule systems. *Computational linguistics* 20(3), 331–378.
- Karttunen, L. (2010). Update on finite state morphology tools. *Ms., Xerox Palo Alto Research Center*.
- Kozen, D. (2001). Automata on guarded strings and applications. Technical report, Cornell University.
- Lewis, D. (1986). *On the plurality of worlds*, Volume 322. Oxford Blackwell.
- McCarthy, J. (1963). Situations, actions, and causal laws. Technical report, Stanford CS.
- Reiter, R. (2001). *Knowledge in Action: Logical foundations for specifying and implementing dynamical systems*. MIT press.
- Rooth, M. (2017). Finite state intensional semantics. In *IWCS 2017-12th International Conference on Computational Semantics-Long papers*.